

Tree-Structured Indexes

SUN YAT-SEN UNIVERSITY

Review

■ Cost Model（代价模型）

□ 只关心磁盘块的IO次数

- 文件由页组成，而每个页包含一组记录
- Record id = <page id, slot #>
- 从随机访问的角度来说，读写一条记录需要一次磁盘IO。

■ 索引技术概述

□ 可以为关系建立索引，都是文件

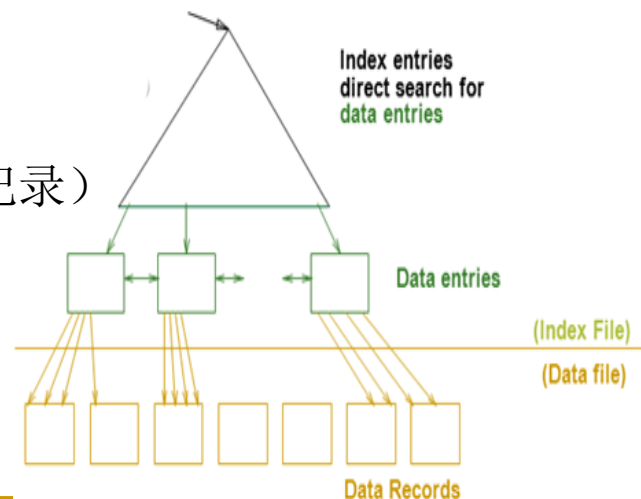
□ 索引文件由两部份组成

1. 数据项部分

- **Data Entry**(数据项) \iff **data record**（数据记录）

2. 引导部份

- 树索引技术
- Hash索引

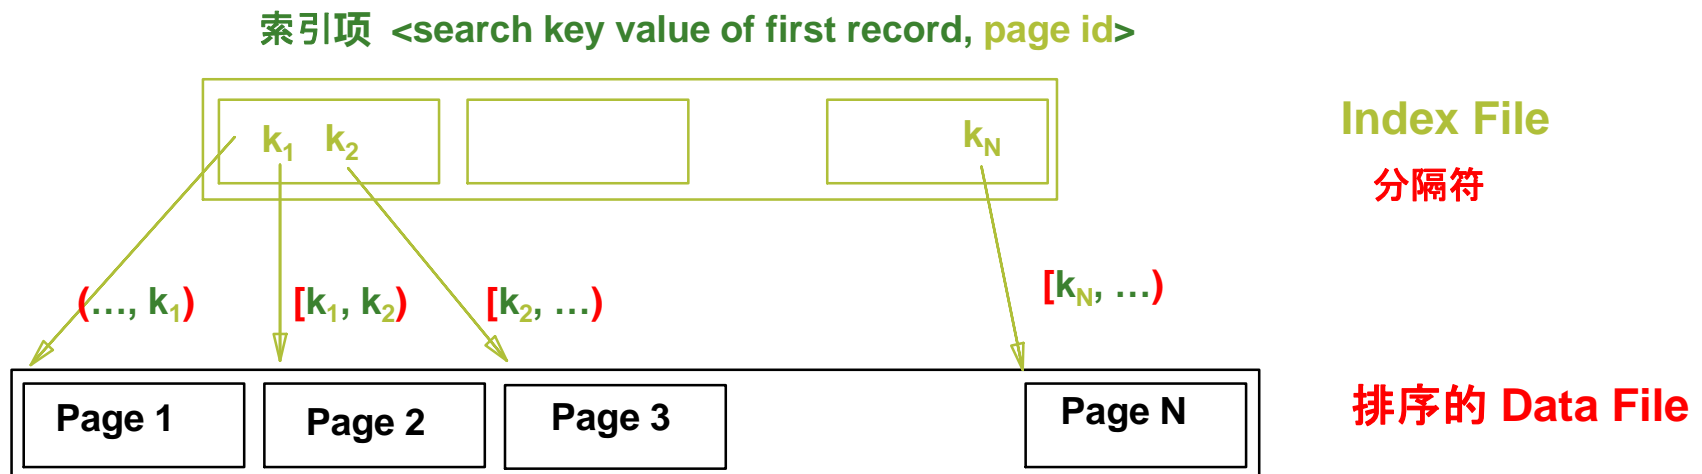


Tree-Structured Indexes: Introduction

- Tree-structured indexing techniques support both *equality selections* and *range selections*.
 - ISAM(索引顺序存取方法): static structure; early index technology.
 - B+ tree: dynamic, adjusts gracefully under inserts and deletes.

Range Searches


- ``Find all students with $gpa > 3.0$ ``
 - ❑ If data is in **sorted file**, do **binary search** to find first such student, then scan to find others.
 - ❑ Cost of **binary search** in a database can be quite high.
 - Why???
- Simple idea: Create an `index' file.




Can do binary search on (smaller) index file!

ISAM

index entry (分隔符)



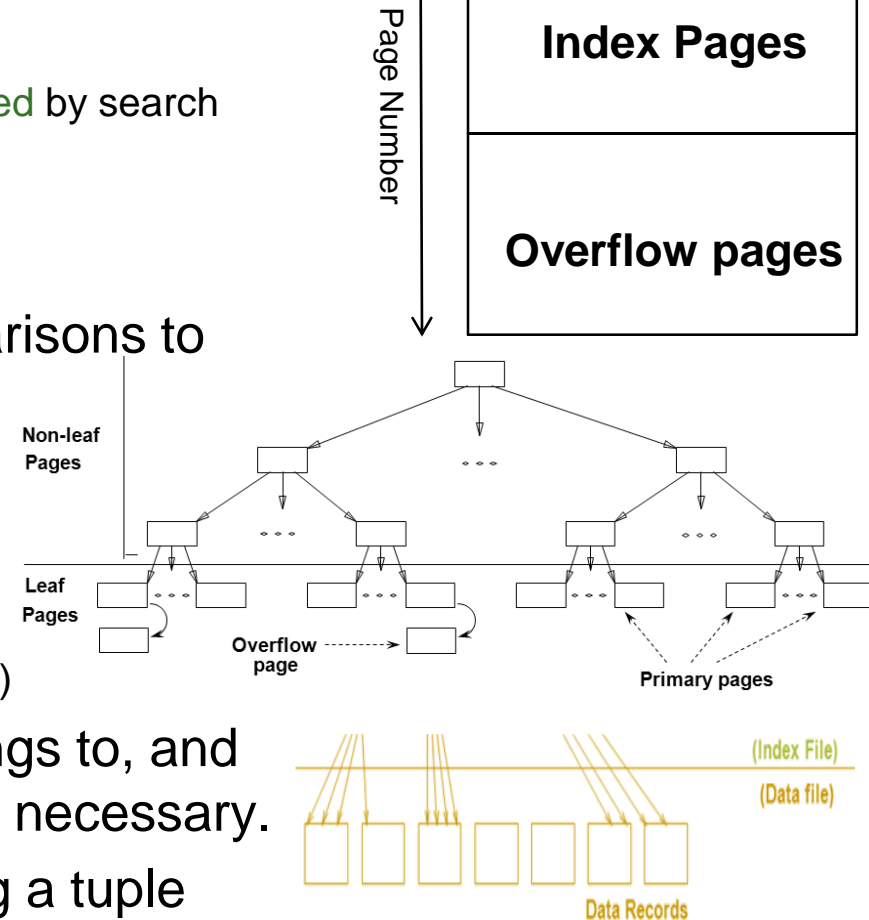
- Index file may still be quite large. But we can apply the idea repeatedly!



✉ Leaf pages contain *data entries*.

ISAM is a STATIC Structure

- **File creation:**
 - Leaf (data) pages allocated **sequentially, sorted** by search key
 - then index pages
 - then overflow pgs.
- **Search:** Start at root; use key comparisons to go to leaf.
- **Cost** = $\log_F N$
 - F = # entries/page (i.e., fanout 扇出)
 - N = # leaf pages
 - no need for 'next-leaf-page' pointers. (Why?)
- **Insert:** Find leaf that data entry belongs to, and put it there. Overflow page(溢出页) if necessary.
- **Delete:** Seek and destroy! If deleting a tuple empties an overflow page, de-allocate it and remove from linked-list.

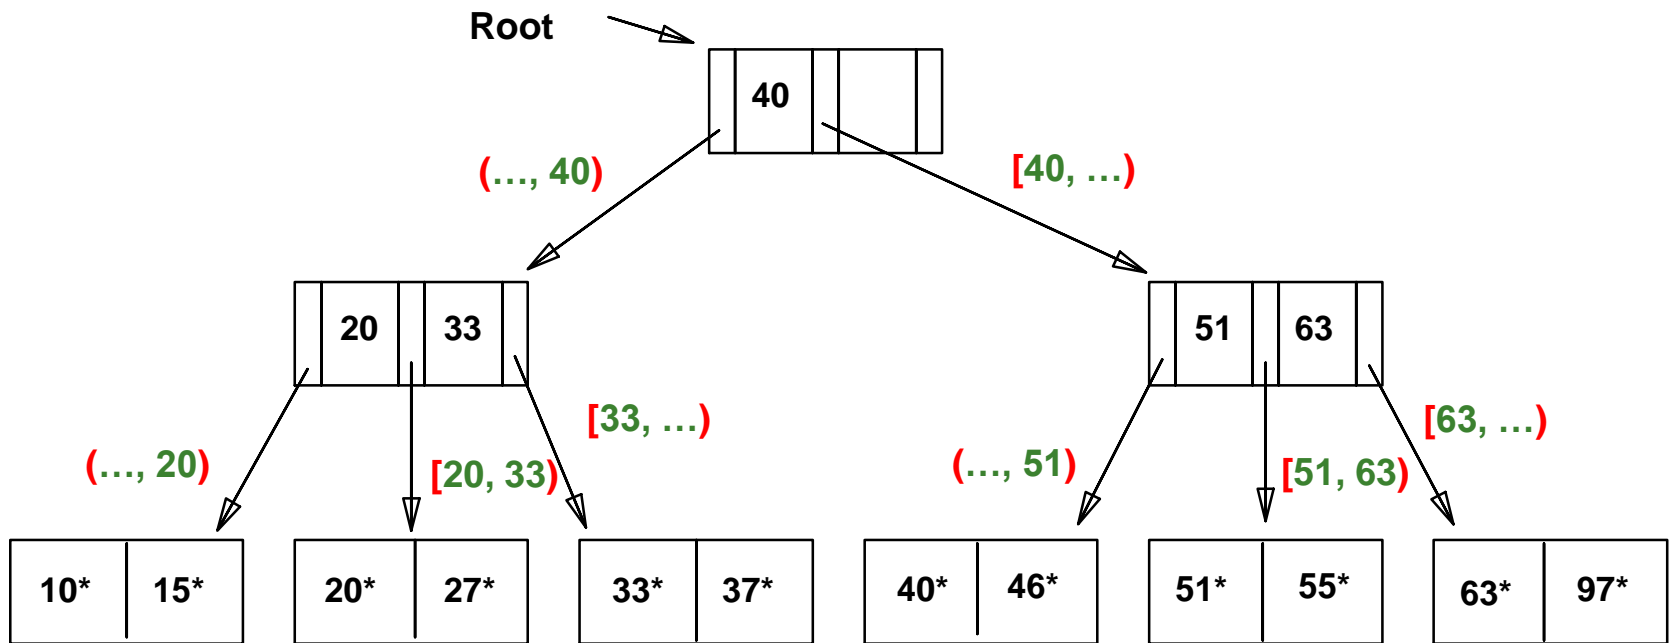


Static tree structure: *inserts/deletes affect only leaf pages.*

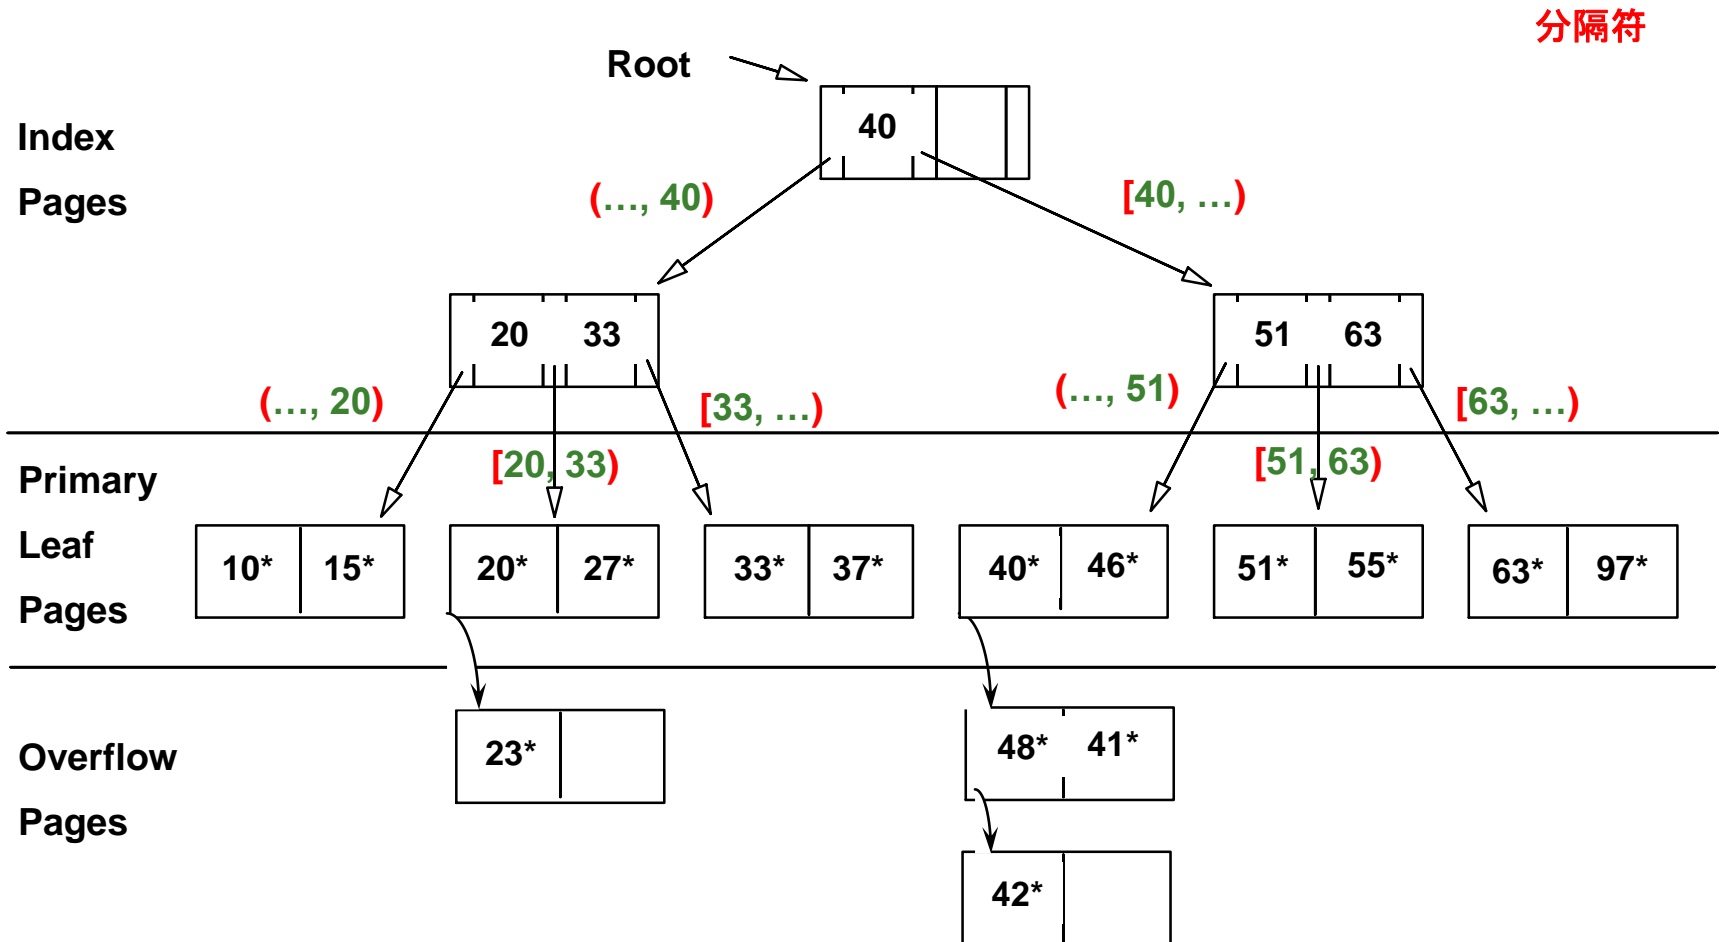
Example ISAM Tree

- Example where each node can hold 2 entries;
- *Index entries*: $\langle \text{search key value, page id} \rangle$, they direct search for data entries *in leaves*.

分隔符

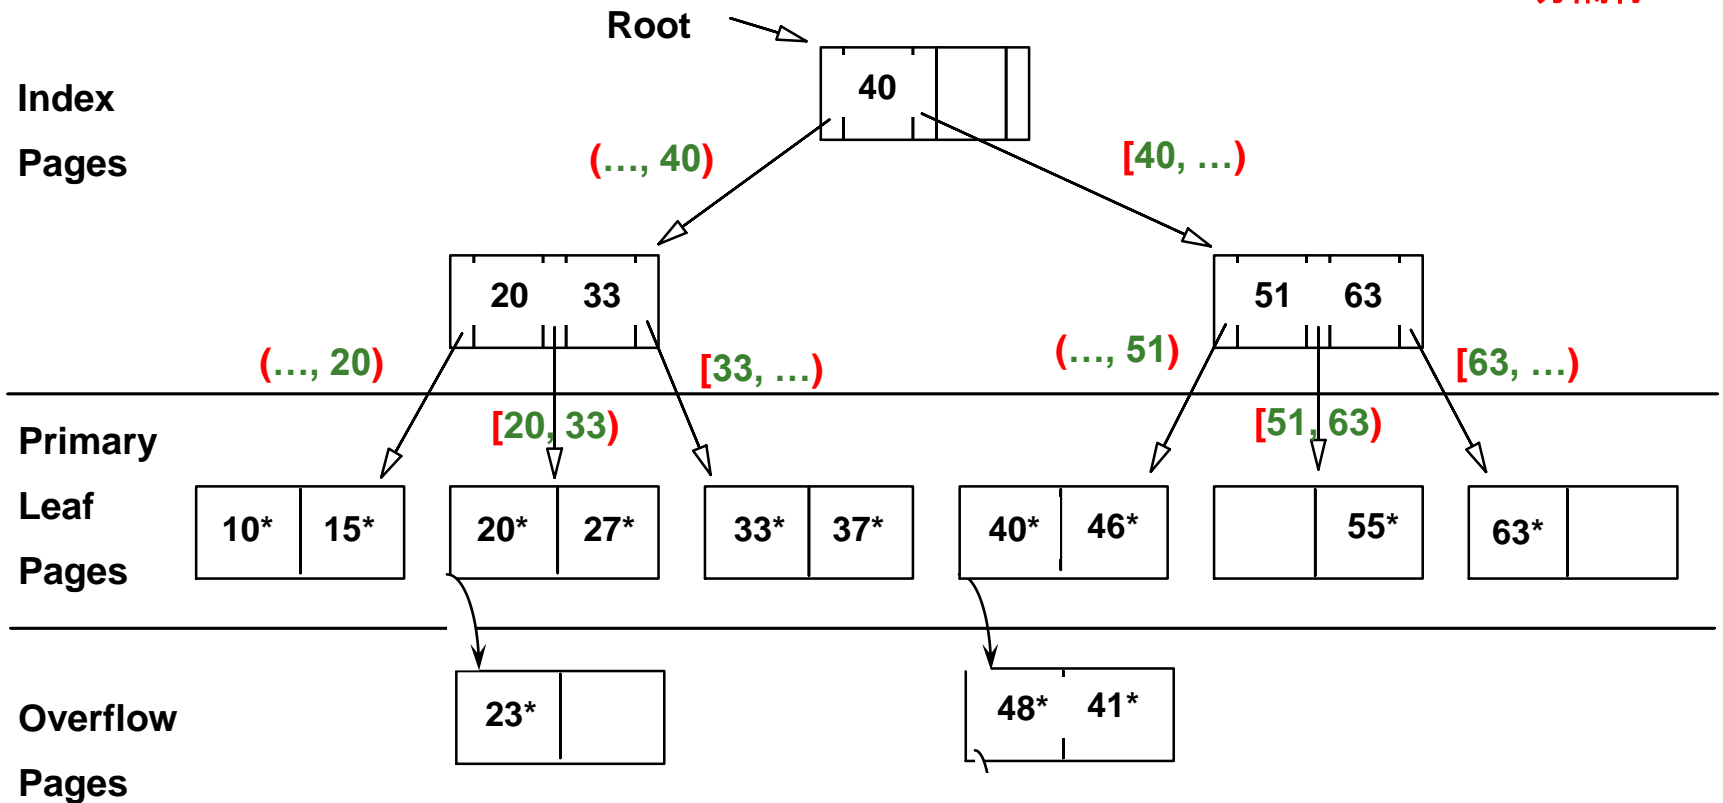


Example: Insert 23^* , 48^* , 41^* , 42^*



... then Deleting 42*, 51*, 97*

分隔符

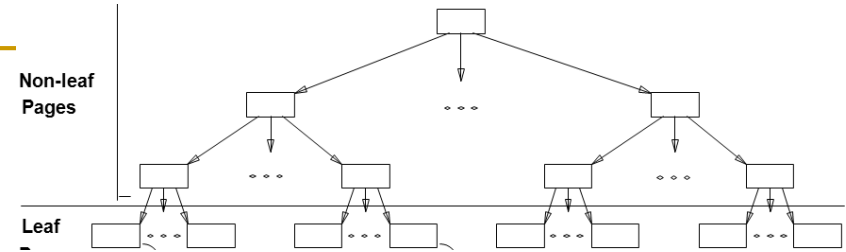


ISAM is a **STATIC** Structure

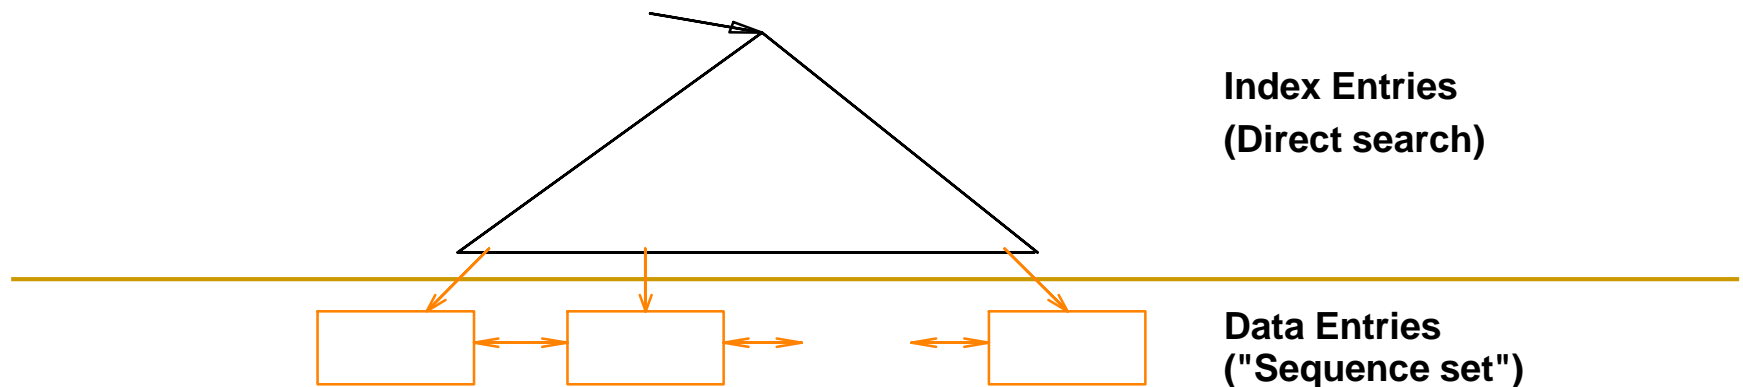


Note that 51 appears in index levels, but not in leaf!

B+ Tree Structure

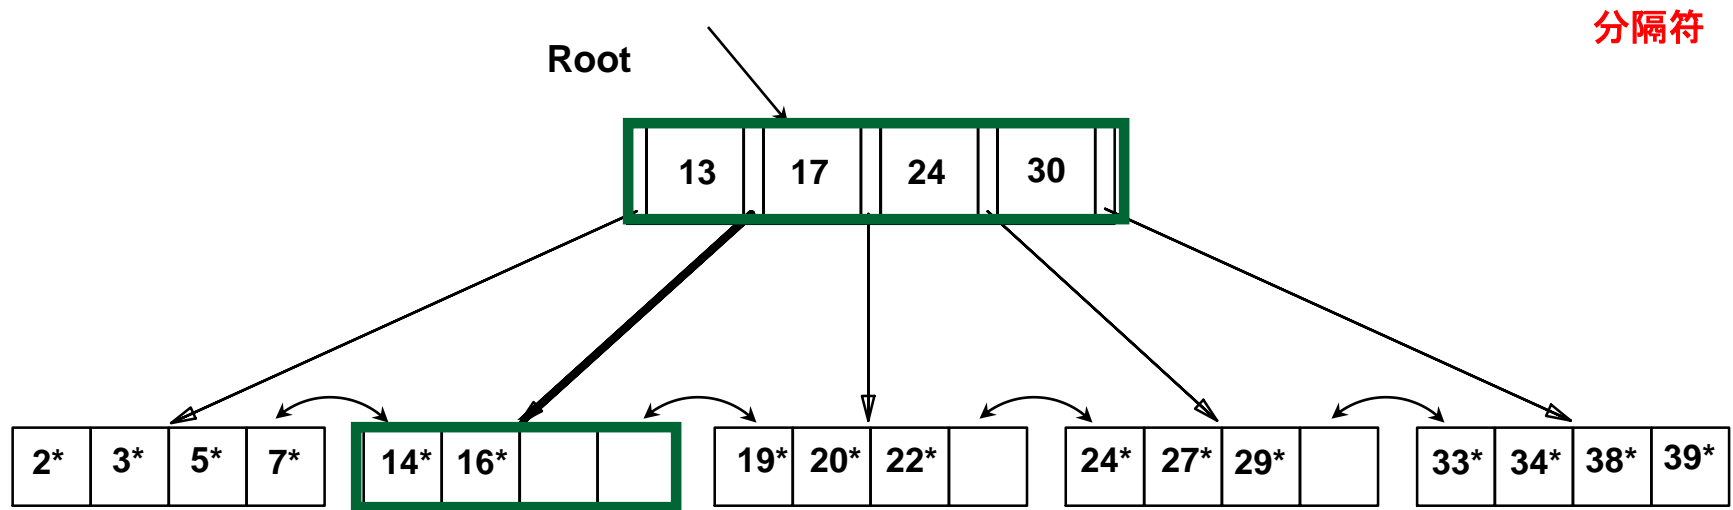


- Each node contains $d \leq m \leq 2d$ entries (index or data)
 - The parameter d is called the **order**(**秩**) of the tree.
 - Each internal node contains m **index entries**: <key, page id>.
 - Each leaf node contains m **data entries**: <key, record **or** record id>
- The ROOT node contains between 1 and $2d$ index entries.
 - It is a leaf or has at least two children.
- Each path from the ROOT to any leaf has the **same length**.
 - 平衡树 -- 树的高度
- Supports equality and range-searches **efficiently**.



B+ Tree Equality Search

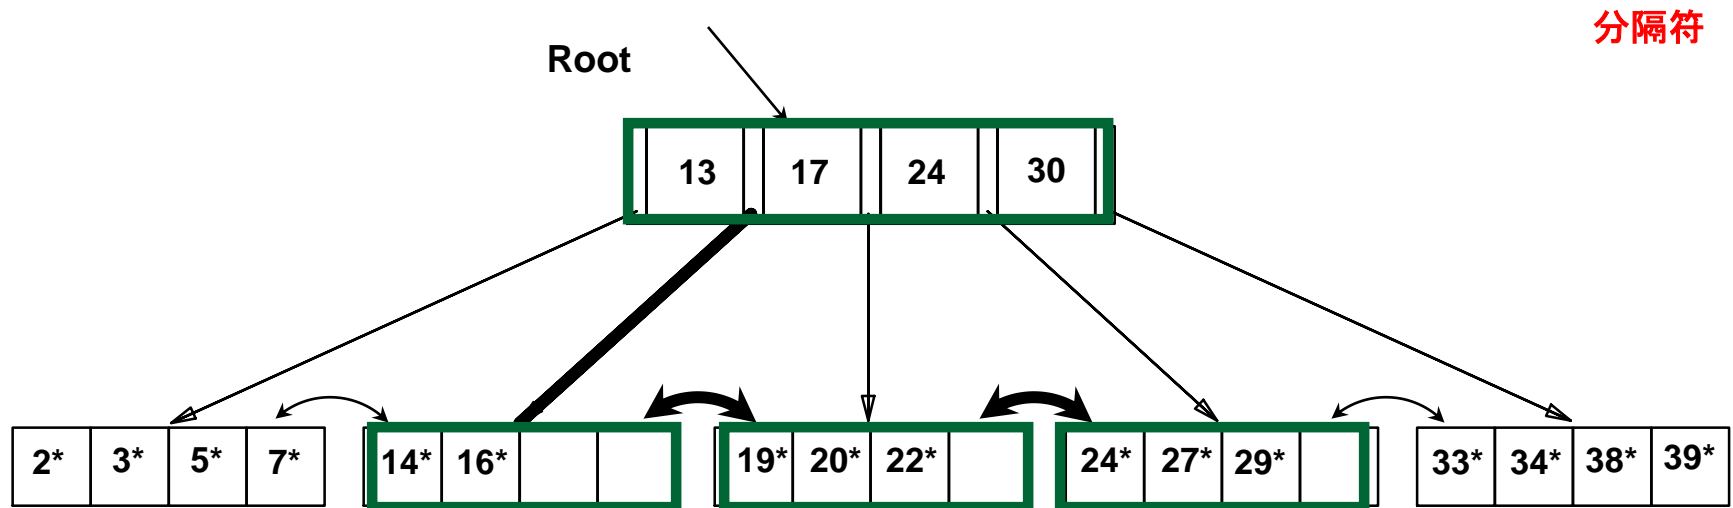
- Search begins at root, and key comparisons direct it to a leaf.
- Search for 15*...



Based on the search for 15, we know it is not in the tree!*

B+ Tree Range Search

- Search all records whose ages are in [15,28].
 - Equality search 15*.
 - Follow sibling **pointers**.



B+ Trees in Practice

$$d \leq \underline{m} \leq 2d$$

- Typical order (秩) : 100. Typical fill-factor: 67%.

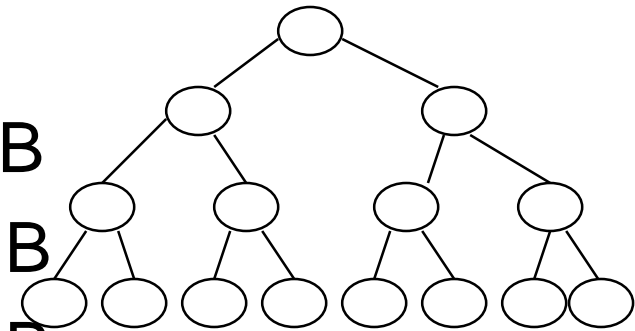
□ average fanout(扇出) = 133 $2d^{2/3}$

□ Level 1 = 1 page = 8 KB

□ Level 2 = 133 pages = 1 MB

□ Level 3 = 17,689 pages = 145 MB

□ Level 4 = 2,352,637 pages = 19 GB



- Can often hold top levels in buffer pool:
 - ❖ With 1 MB buffer, can locate one record in 19 GB (or 0.3 billion records) in two I/Os!